

Pakshya AI

Additional Technical & Tricky Reviewer Questions

Additional Technical Questions

1. How is the image uploaded from React to FastAPI, and what data format is used?
2. How do you handle multiple images belonging to the same inspection?
3. How do you maintain the relationship between an image and its extracted declarations?
4. Where are inspection images and extracted results stored?
5. How do you prevent unauthorized modification of an inspection record?
6. How would JWT token expiry be handled during an inspection?
7. What happens if a JWT is stolen or compromised?
8. How do you separate authentication from authorization in your backend?
9. How are different inspector and senior-officer permissions represented in the system?
10. How would you design your API endpoints for the inspection workflow?
11. How do you validate incoming data before sending it to the compliance engine?
12. What happens if the OCR service fails completely?
13. What happens if only one OCR engine is available?
14. How do you decide which OCR output should be sent to the declaration-analysis stage?
15. How do you handle rotated or upside-down text?
16. How do you handle text appearing at different orientations on the same package?
17. How does your system deal with OCR detecting text that is not a legal declaration?
18. How do you distinguish MRP from other prices printed on the package?
19. How do you distinguish net quantity from other numbers appearing on the package?
20. How do you handle different units such as g, kg, ml and L during normalization?
21. How do you normalize different ways of writing the same declaration?
22. How does the system prevent one extracted field from being incorrectly mapped to another field?
23. What happens when a required declaration cannot be extracted?
24. How do you represent an uncertain declaration internally?
25. How does the rule engine consume the normalized data?

More Tricky Questions

1. What happens if the package contains text in two or more languages and the declarations contradict each other?

2. What happens if OCR extracts a valid-looking MRP but the currency symbol is missing?
3. How would you detect whether ■100 is actually MRP or simply a promotional price?
4. If the same declaration appears multiple times on a package, which occurrence does your system validate?
5. How do you handle handwritten markings or manually added information on a package?
6. How would you handle transparent packaging where text from the opposite side is visible?
7. How do you prevent background text or advertisements from being interpreted as mandatory declarations?
8. If a barcode points to a different product than the visible package, how is that discrepancy handled?
9. What happens when the barcode cannot be scanned?
10. What happens when the barcode is valid but the online product information has changed?
11. How do you distinguish an actual product change from a simple change in image quality between inspections?
12. How do you ensure image hashing does not incorrectly treat two different package versions as identical?
13. If two inspectors inspect the same package and obtain different results, how is this represented in the inspection history?
14. How would you detect repeated violations when the same product has a slightly different name or packaging?
15. What exactly makes your system's evidence "explainable" rather than simply showing an AI prediction?

Very Difficult Reviewer Questions

1. Which component of your architecture is responsible for deciding compliance, and which component is responsible only for extracting information?
2. If your OCR is 95% accurate, does that mean your overall compliance system is 95% accurate? Why or why not?
3. What is the difference between OCR accuracy, field-extraction accuracy and final compliance-decision accuracy?
4. Which evaluation metrics would you use separately for sticker detection, OCR, declaration extraction and rule validation?
5. How would you construct a test dataset for evaluating the complete pipeline?
6. How would you test the system against packages containing deliberately misleading or ambiguous declarations?
7. How would you test whether your rule engine is applying the correct legal requirement rather than simply producing a technically valid result?
8. How do you reproduce and debug a wrong compliance decision made several months earlier?

- 9.** If the AI and rule engine produce a technically correct result but the officer disagrees, how do you determine whether the AI or the legal interpretation was wrong?
- 10.** What is the single most critical point of failure in your complete pipeline, and how would you detect that failure before presenting a result to an officer?